



CS & IT ENGINEERING



Operating Systems-Revision

Process Concepts

Lecture No. - 02

By- Dr. Khaleel Khan
Sir



Recap of Previous Lecture



Topic

Introduction & Background



Topics to be Covered



Topic

Process Concepts

Topic

Threads & Multithreading

Program / Job

(.exe)

Instns

Data

Ex: Load AC, R₁
Add R₁, R₂

Static

Dynamic

→ Fixed &
Known Size

→ Unknown
Variable Size

→ @ Load time
b/f R.T

→ @ R.T alloc.

Process / Task / App



→ Prog in Execution

→ Instance of a Prog

→ Active entity

→ Locus of control

→ unit of cpu utiliz

→ Animated
Spirit



Topic : Process Concept

- An operating system executes a variety of programs that run as a process.
- Process – a program in execution; process execution must progress in sequential fashion. No parallel execution of instructions of a single process
- Multiple parts
 - The program code, also called text section
 - Current activity including program counter, processor registers
 - Stack containing temporary data
 - Function parameters, return addresses, local variables
 - Data section containing global variables
 - Heap containing memory dynamically allocated during run time



Topic : Process Concept (Cont.)

- Program is passive entity stored on disk (executable file); process is active
 - [Program becomes process when an executable file is loaded into memory]
- Execution of program started via GUI mouse clicks, command line entry of its name, etc.
- One program can be several processes (fork() system call)
 - Consider multiple users executing the same program



Topic : Process in Memory

Developer's view of
Process

A.D.T

< Defn, Repr ; operations,
Structure Attributes >

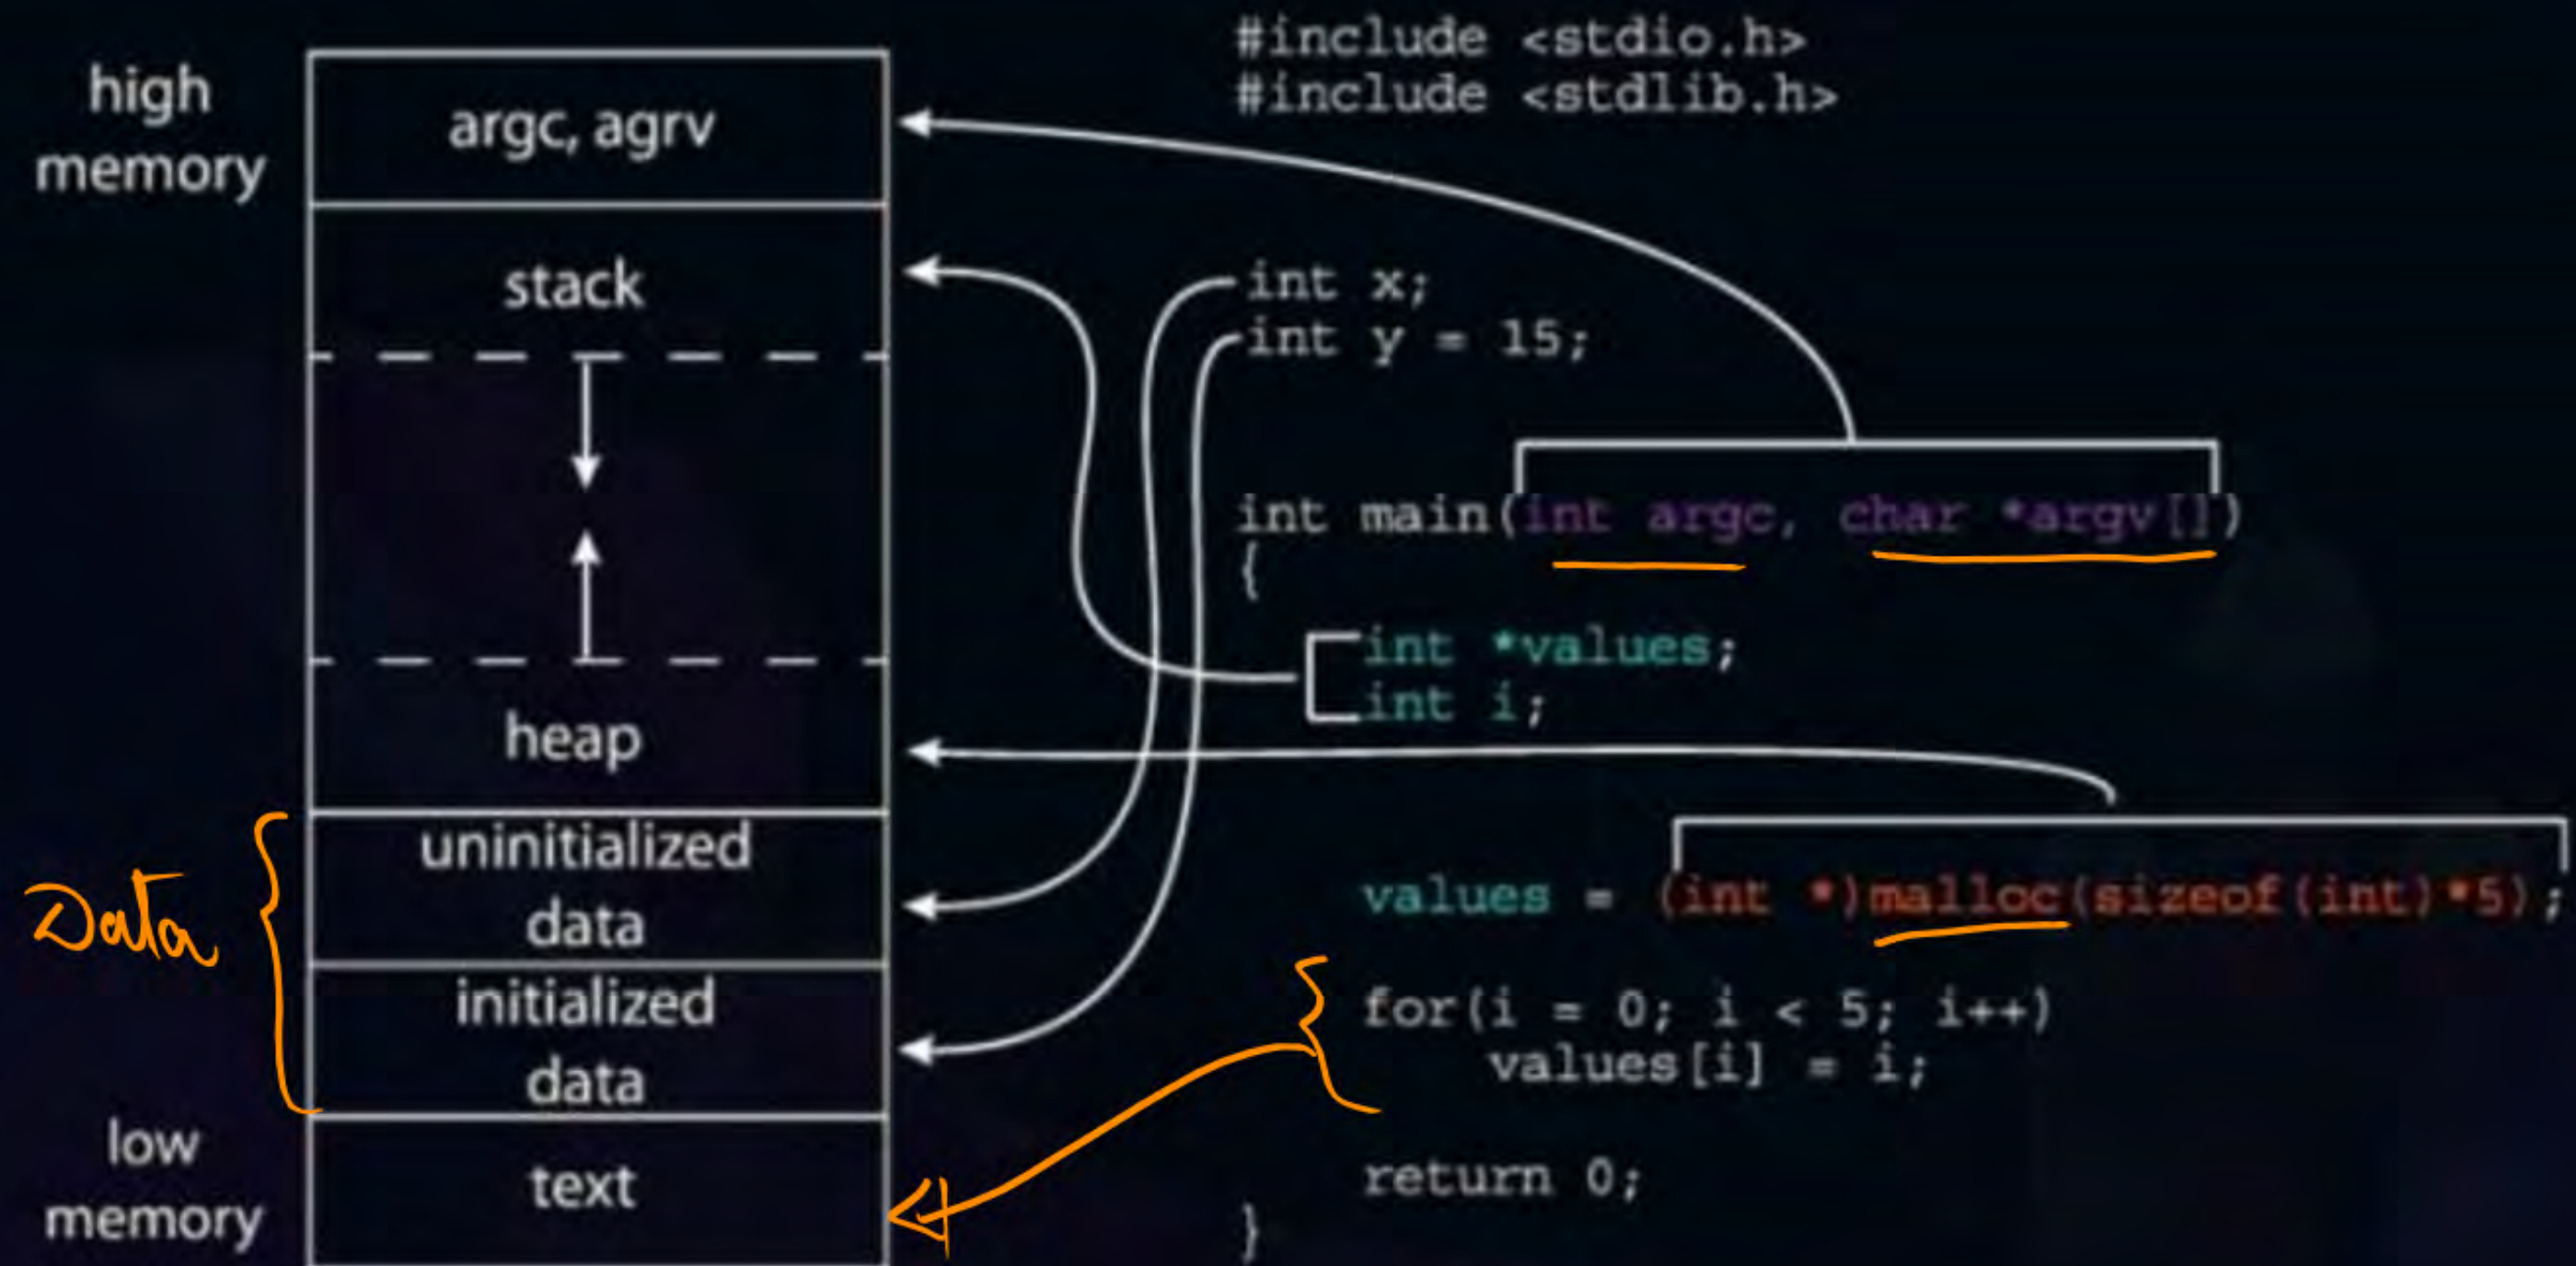


Dynamic Alloc.

Code / Instrns



Topic : Memory Layout of a C Program





Process Attributes:

→ Pid, Ppid, gid,
Pc; G.P.R
mem. limits
Priority; State



→ Process Context
Environment



OS 0.8

K.M by
OS processes

P.C.B (Process Control Block)
Process Descriptor



Topic : Process State

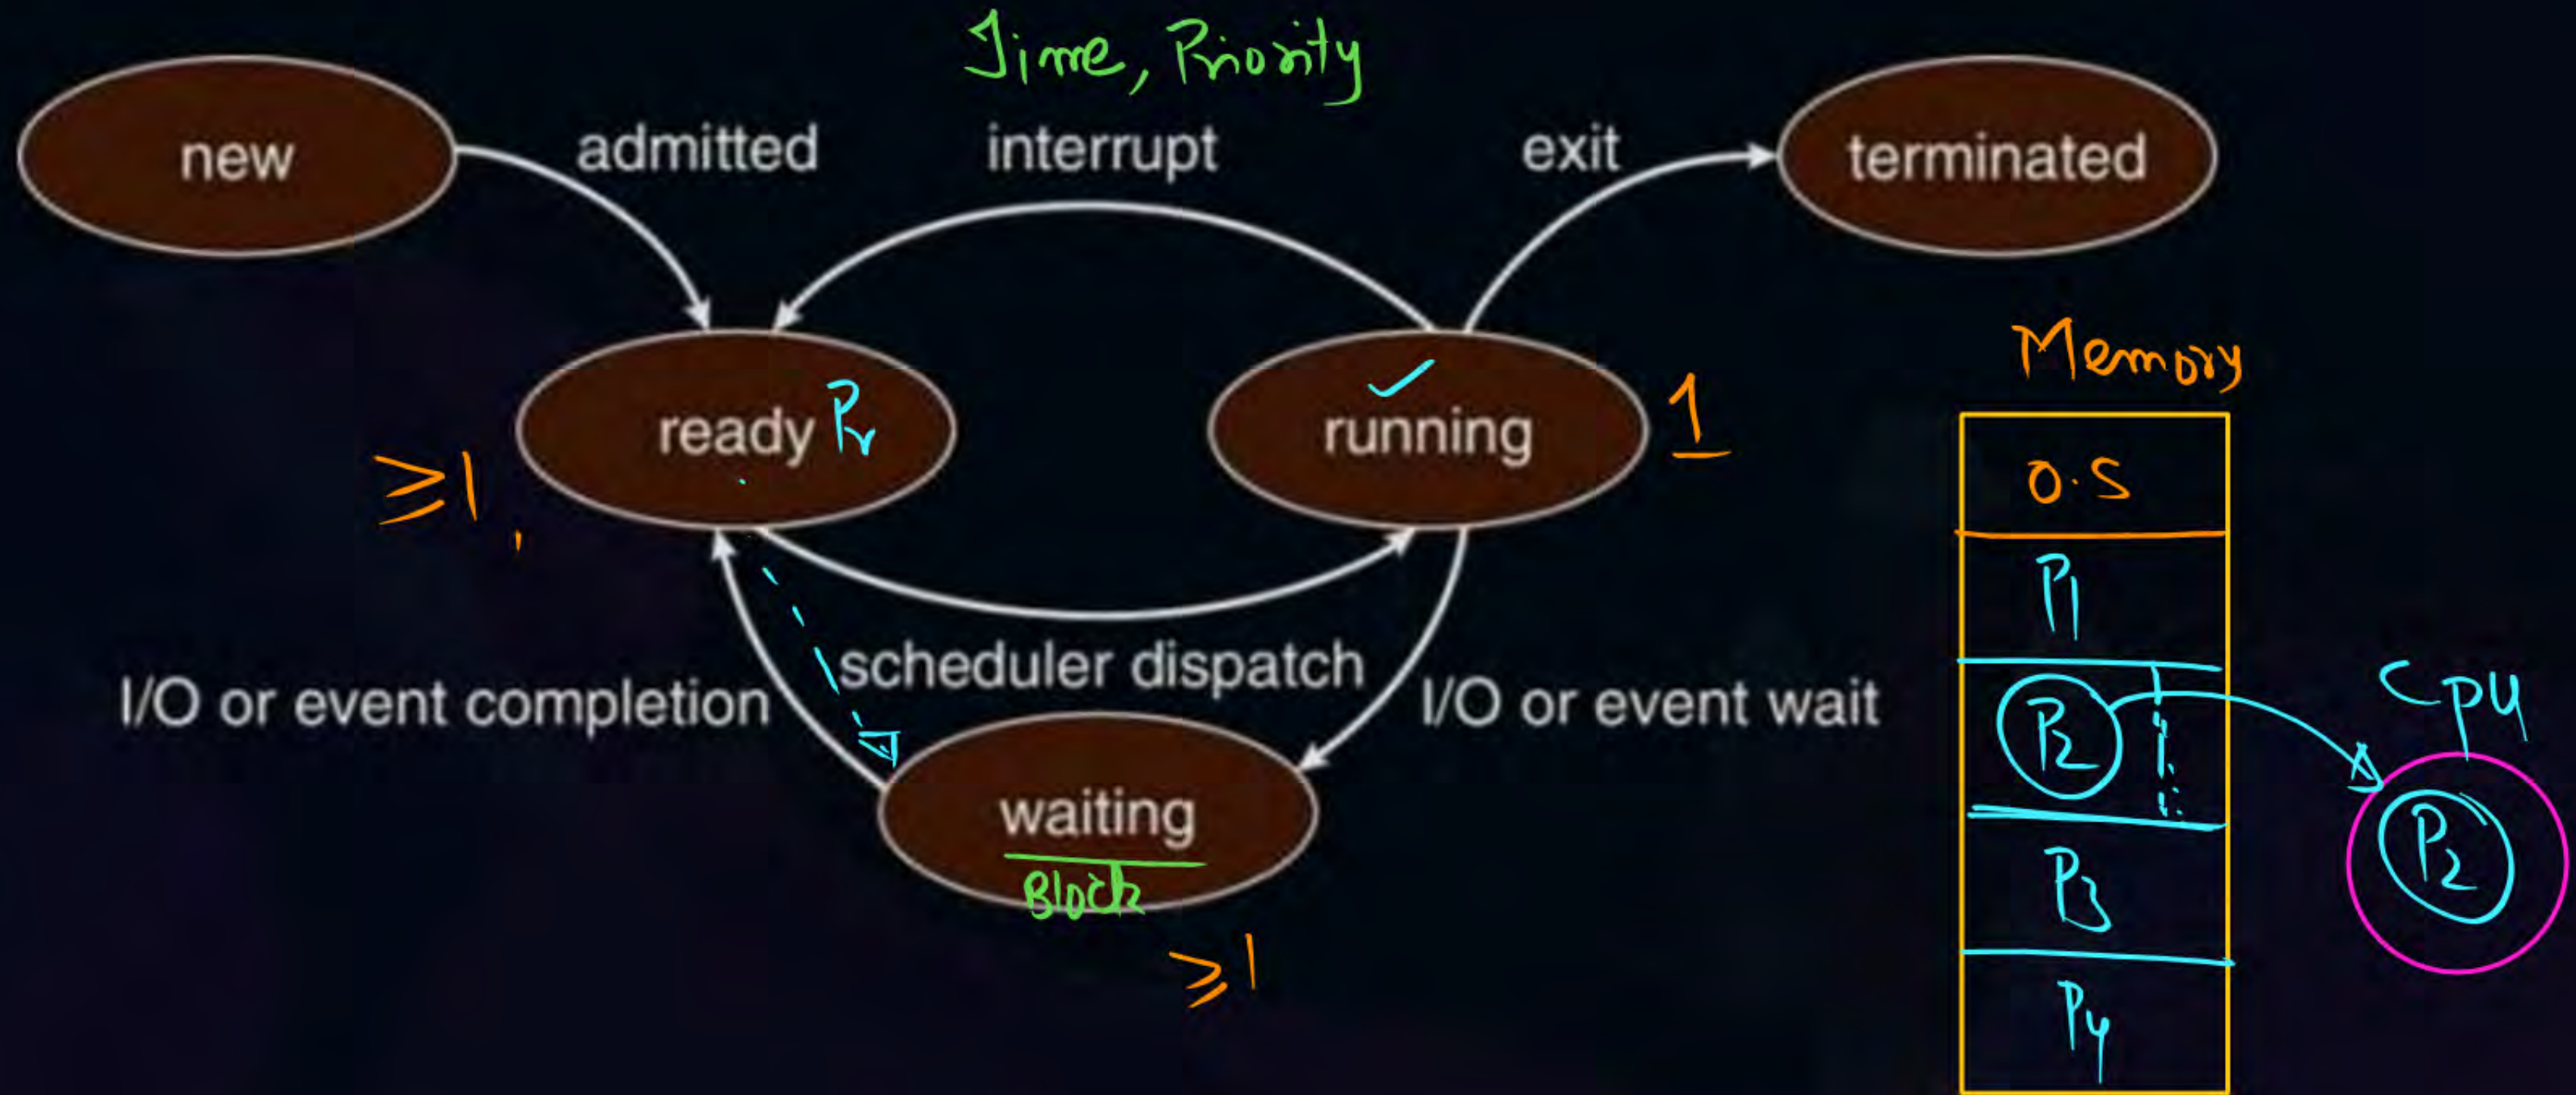
- As a process executes, it changes state
 - New: The process is being created
 - Running: Instructions are being executed *on CPU*
 - Waiting: The process is waiting for some event to occur (*IO, Interrupt (event)*)
 - Ready: The process is waiting to be assigned to a processor (*Ready Q*)
 - Terminated: The process has finished execution

→ Block



Topic : Diagram of Process State

$$\underline{M.R} < \begin{matrix} N.R \\ R \end{matrix}$$





Topic : Process Control Block (PCB)

Information associated with each process(also called task control block)

- Process state – running, waiting, etc.
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files

process state
process number
program counter
registers
memory limits
list of open files
...



Topic : Process Scheduling

- Process scheduler selects among available processes for next execution on CPU core
- Goal -- Maximize CPU use, quickly switch processes onto CPU core
- Maintains scheduling queues of processes
- Ready queue – set of all processes residing in main memory, ready and waiting to execute
- Wait queues – set of processes waiting for an event (i.e., I/O)
- Processes migrate among the various queues

Scheduling Q's & Queuing diagram

In Memory

on disk

Job-Q/IP-Q

Suspend-Q

Ready
Q

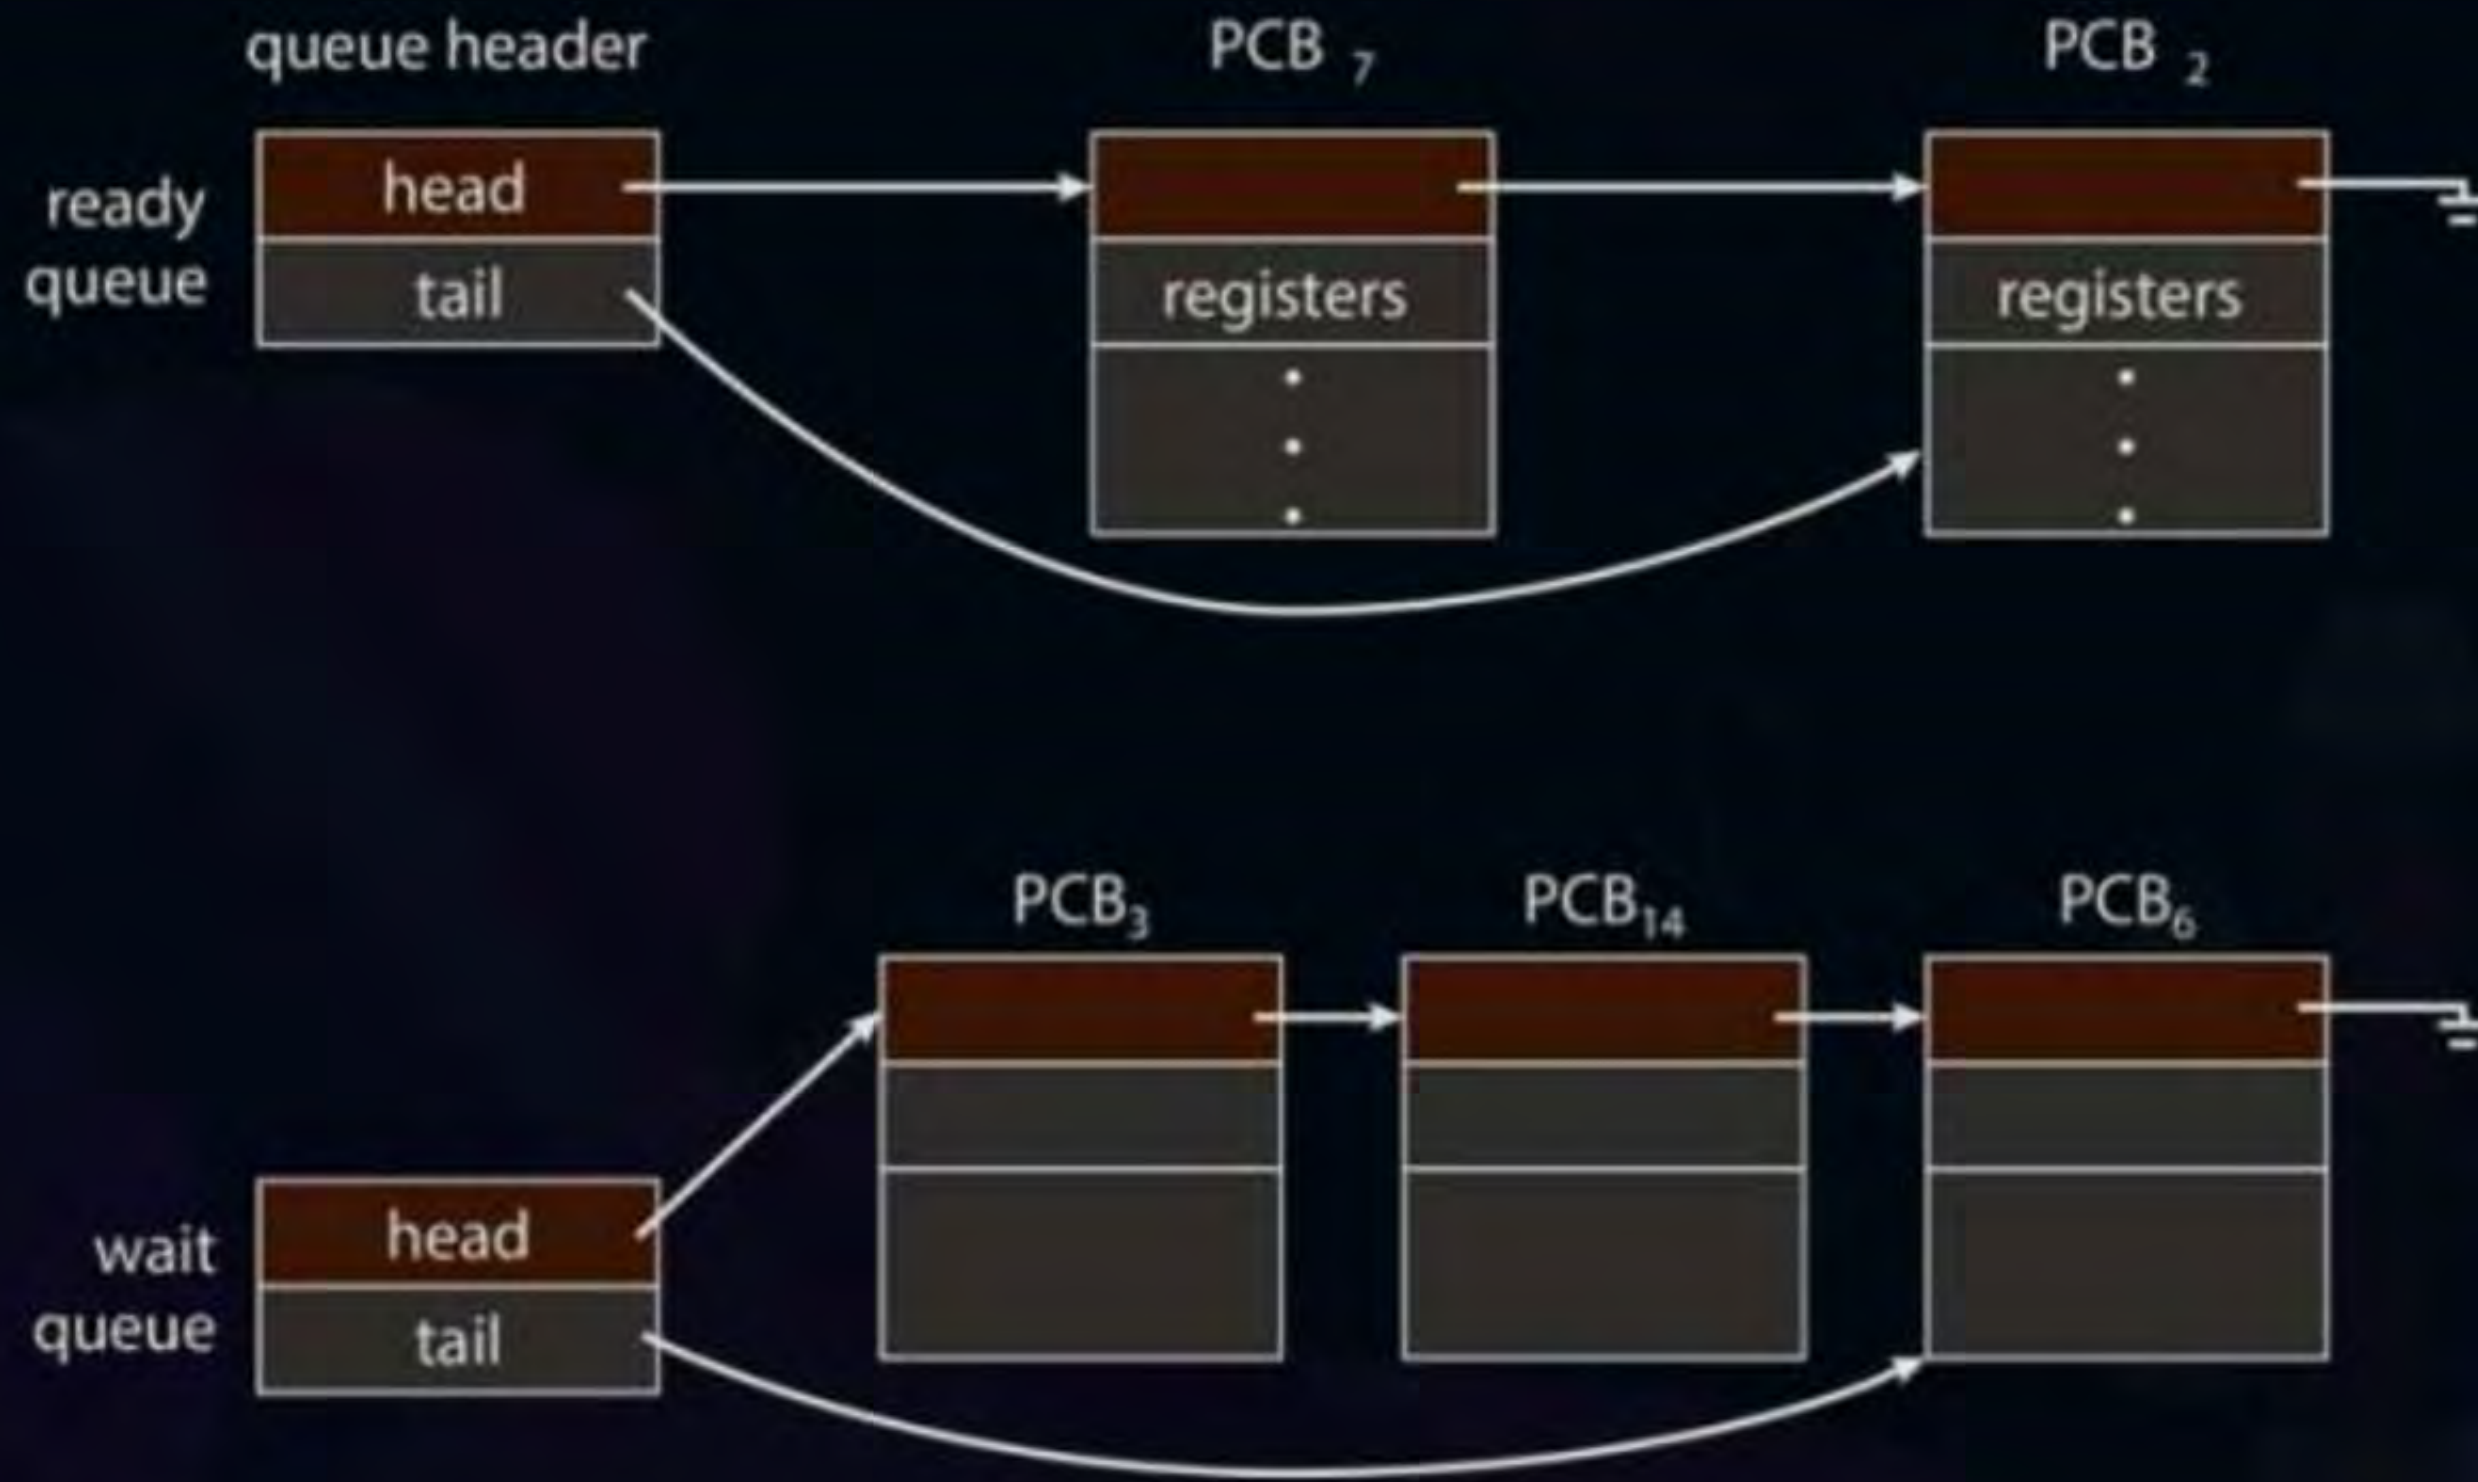
Block
Q
<Io devices>



Linked list

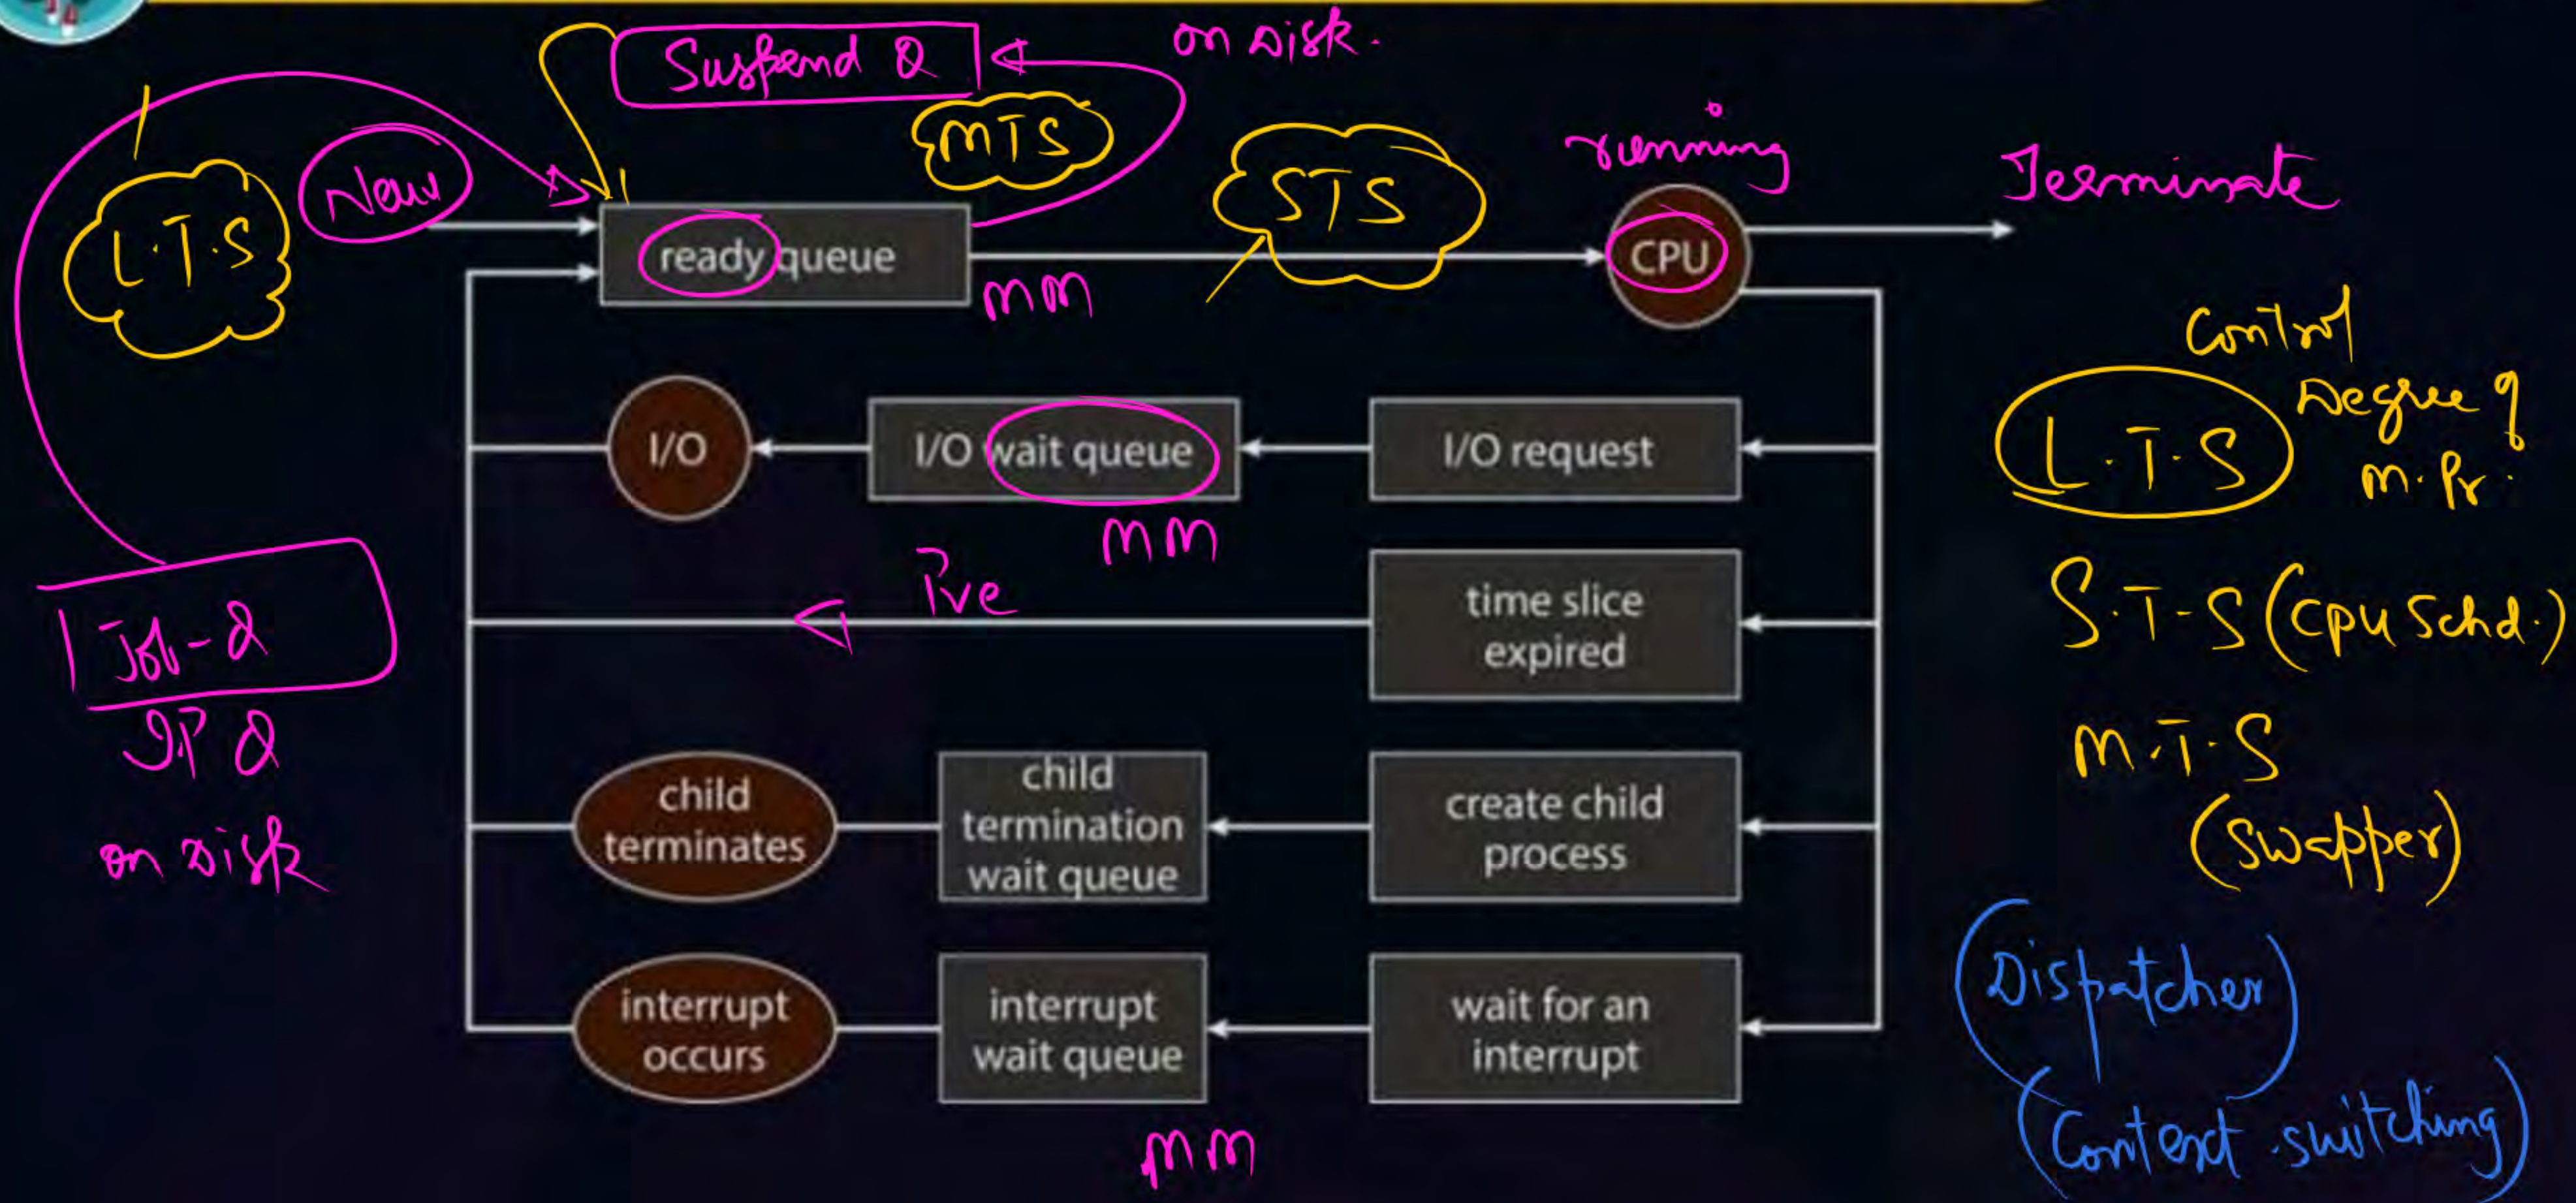


Topic : Ready and Wait Queues





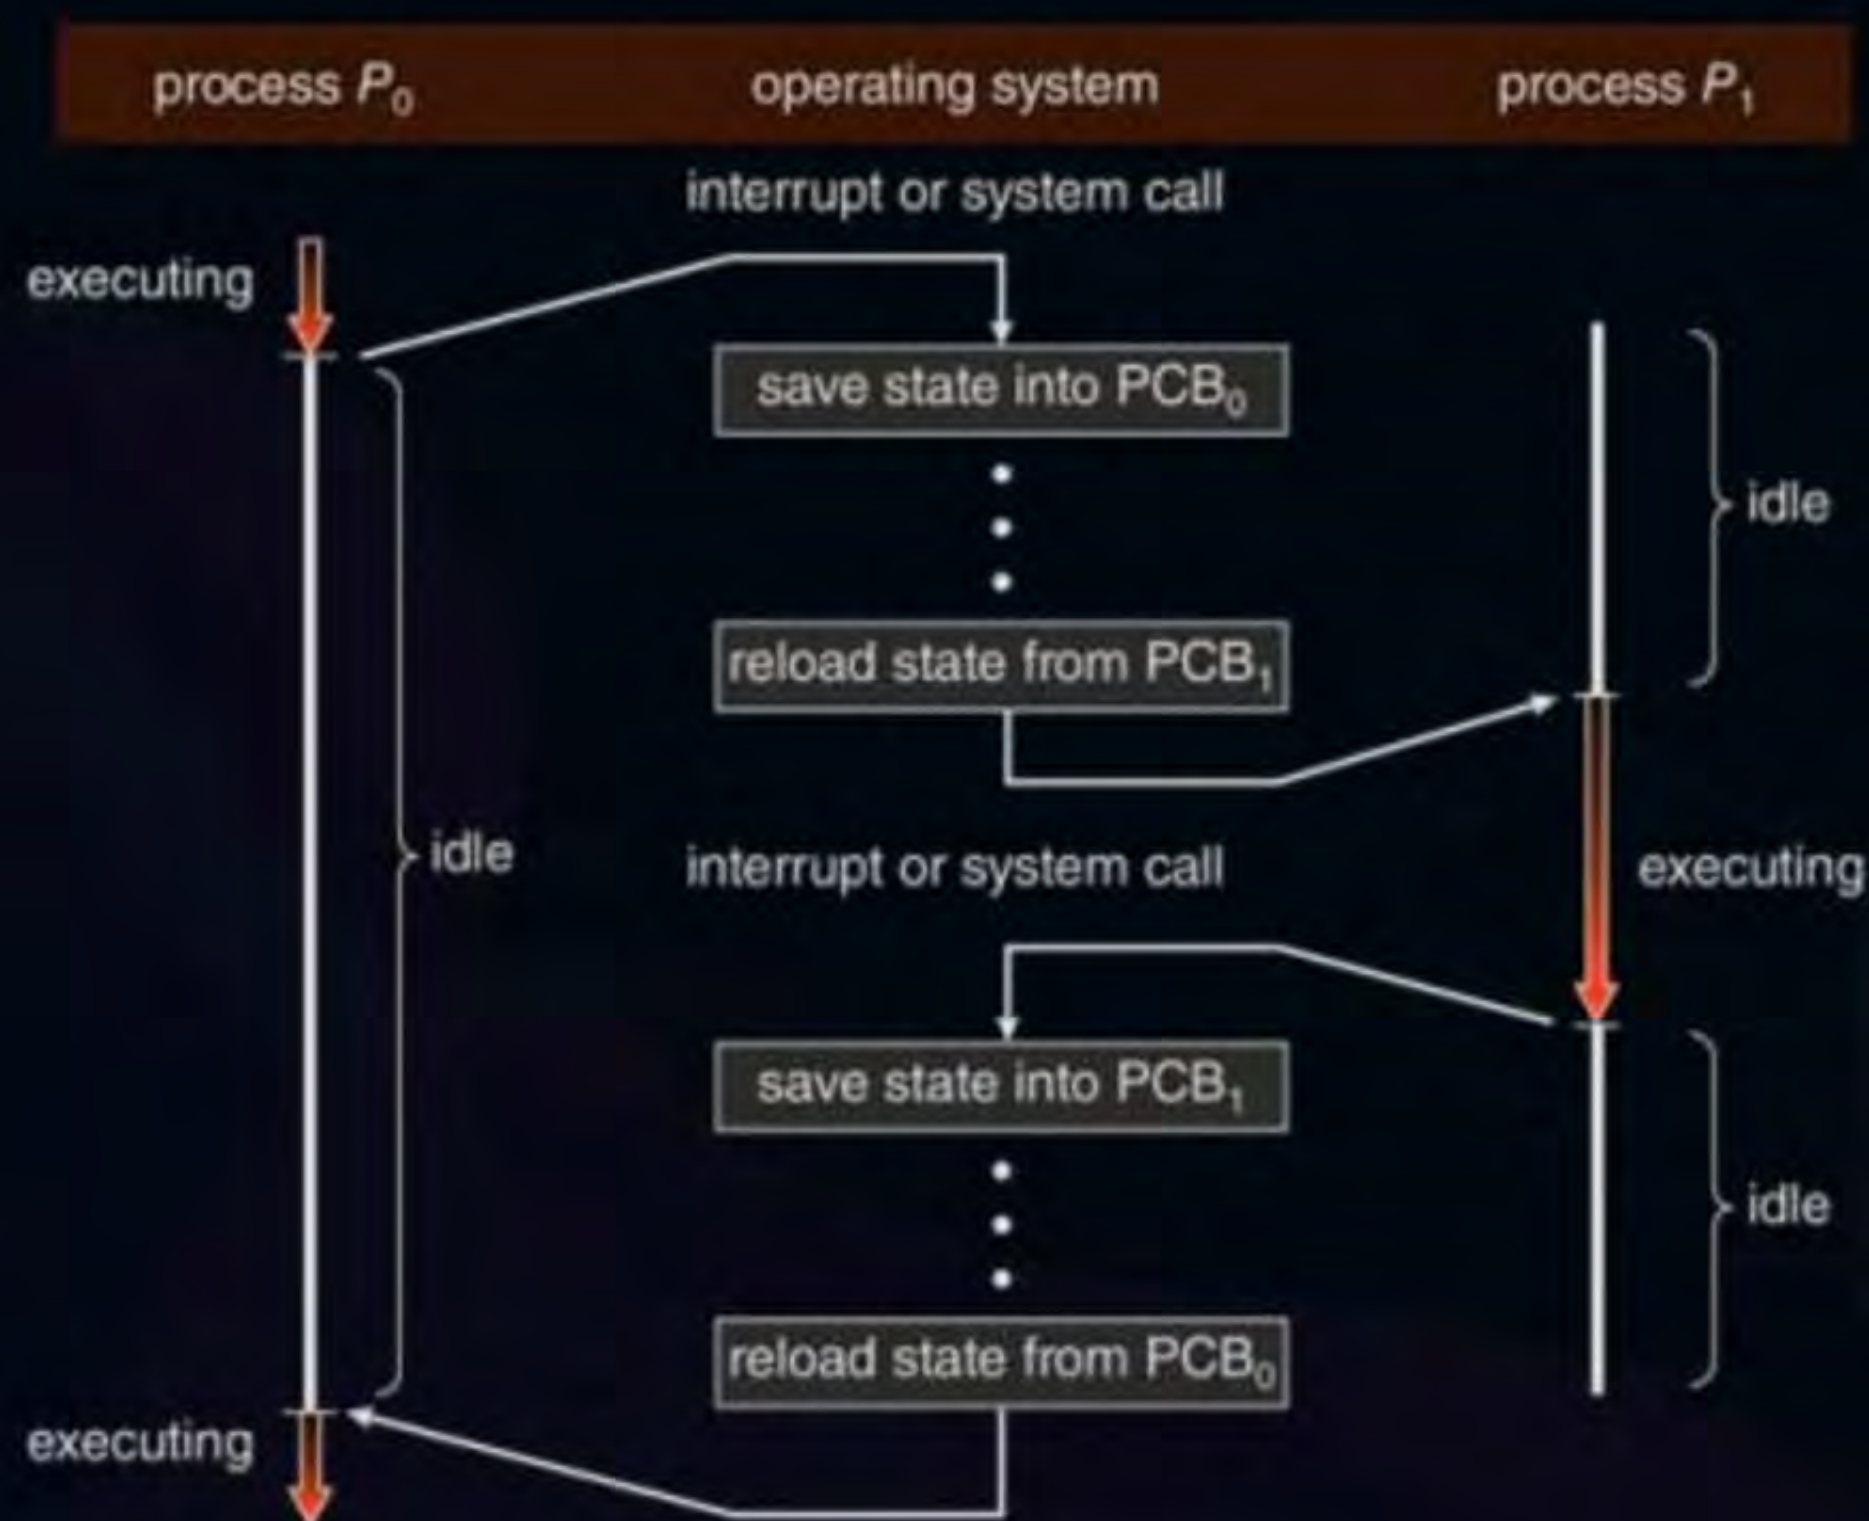
Topic : Representation of Process Scheduling





Topic : CPU Switch From Process to Process

- A context switch occurs when the CPU switches from one process to another.





Topic : Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a context switch
- Context of a process represented in the PCB
- Context-switch time is pure overhead; the system does no useful work while switching
- The more complex the OS and the PCB \Rightarrow the longer the context switch
- Time dependent on hardware support
- Some hardware provides multiple sets of registers per CPU \Rightarrow multiple contexts loaded at once



Topic : Operations on Processes

- System must provide mechanisms for:
- Process creation : *fork*
- Process termination : *exit*



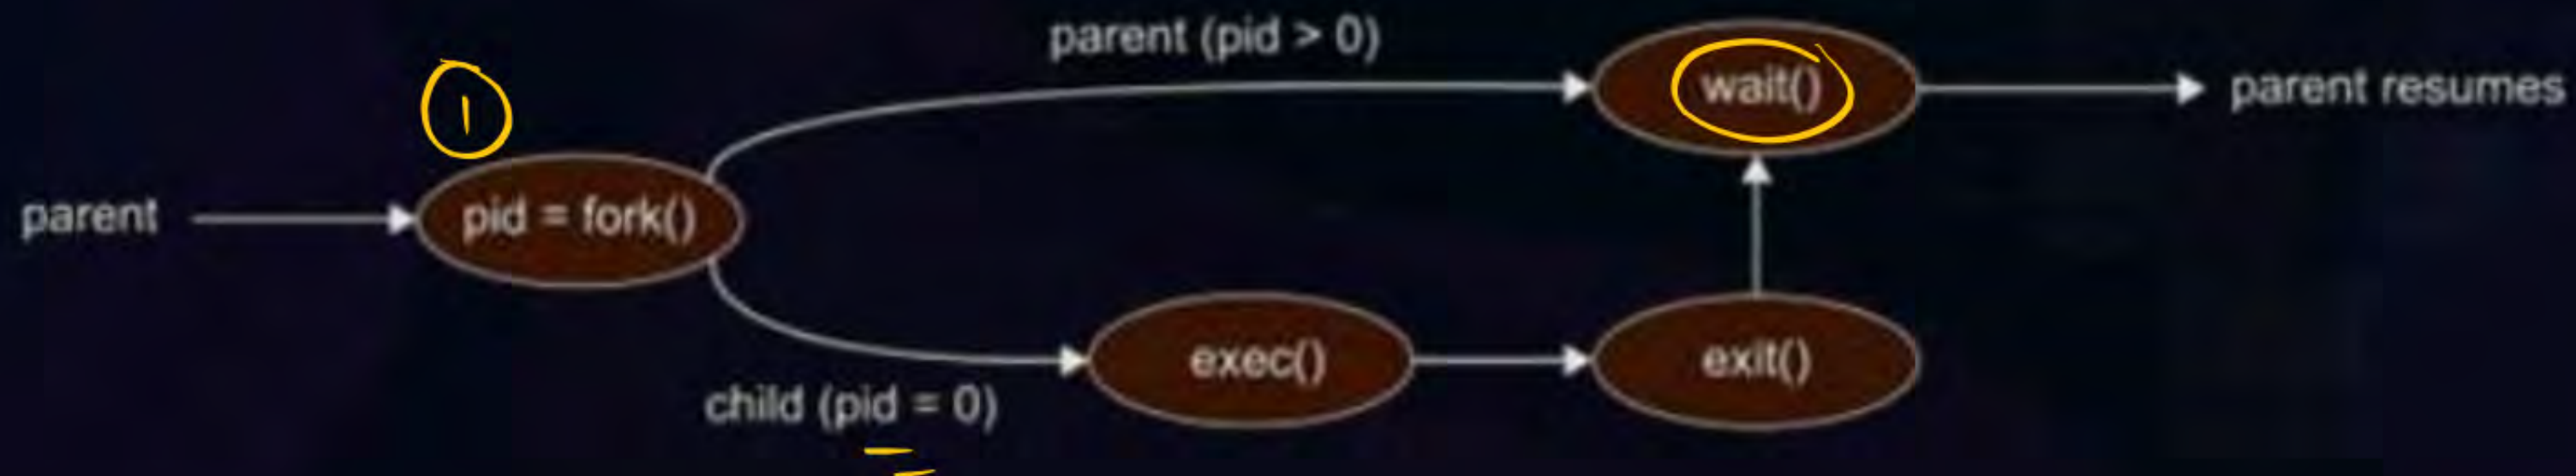
Topic : Process Creation

- Parent process create children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate



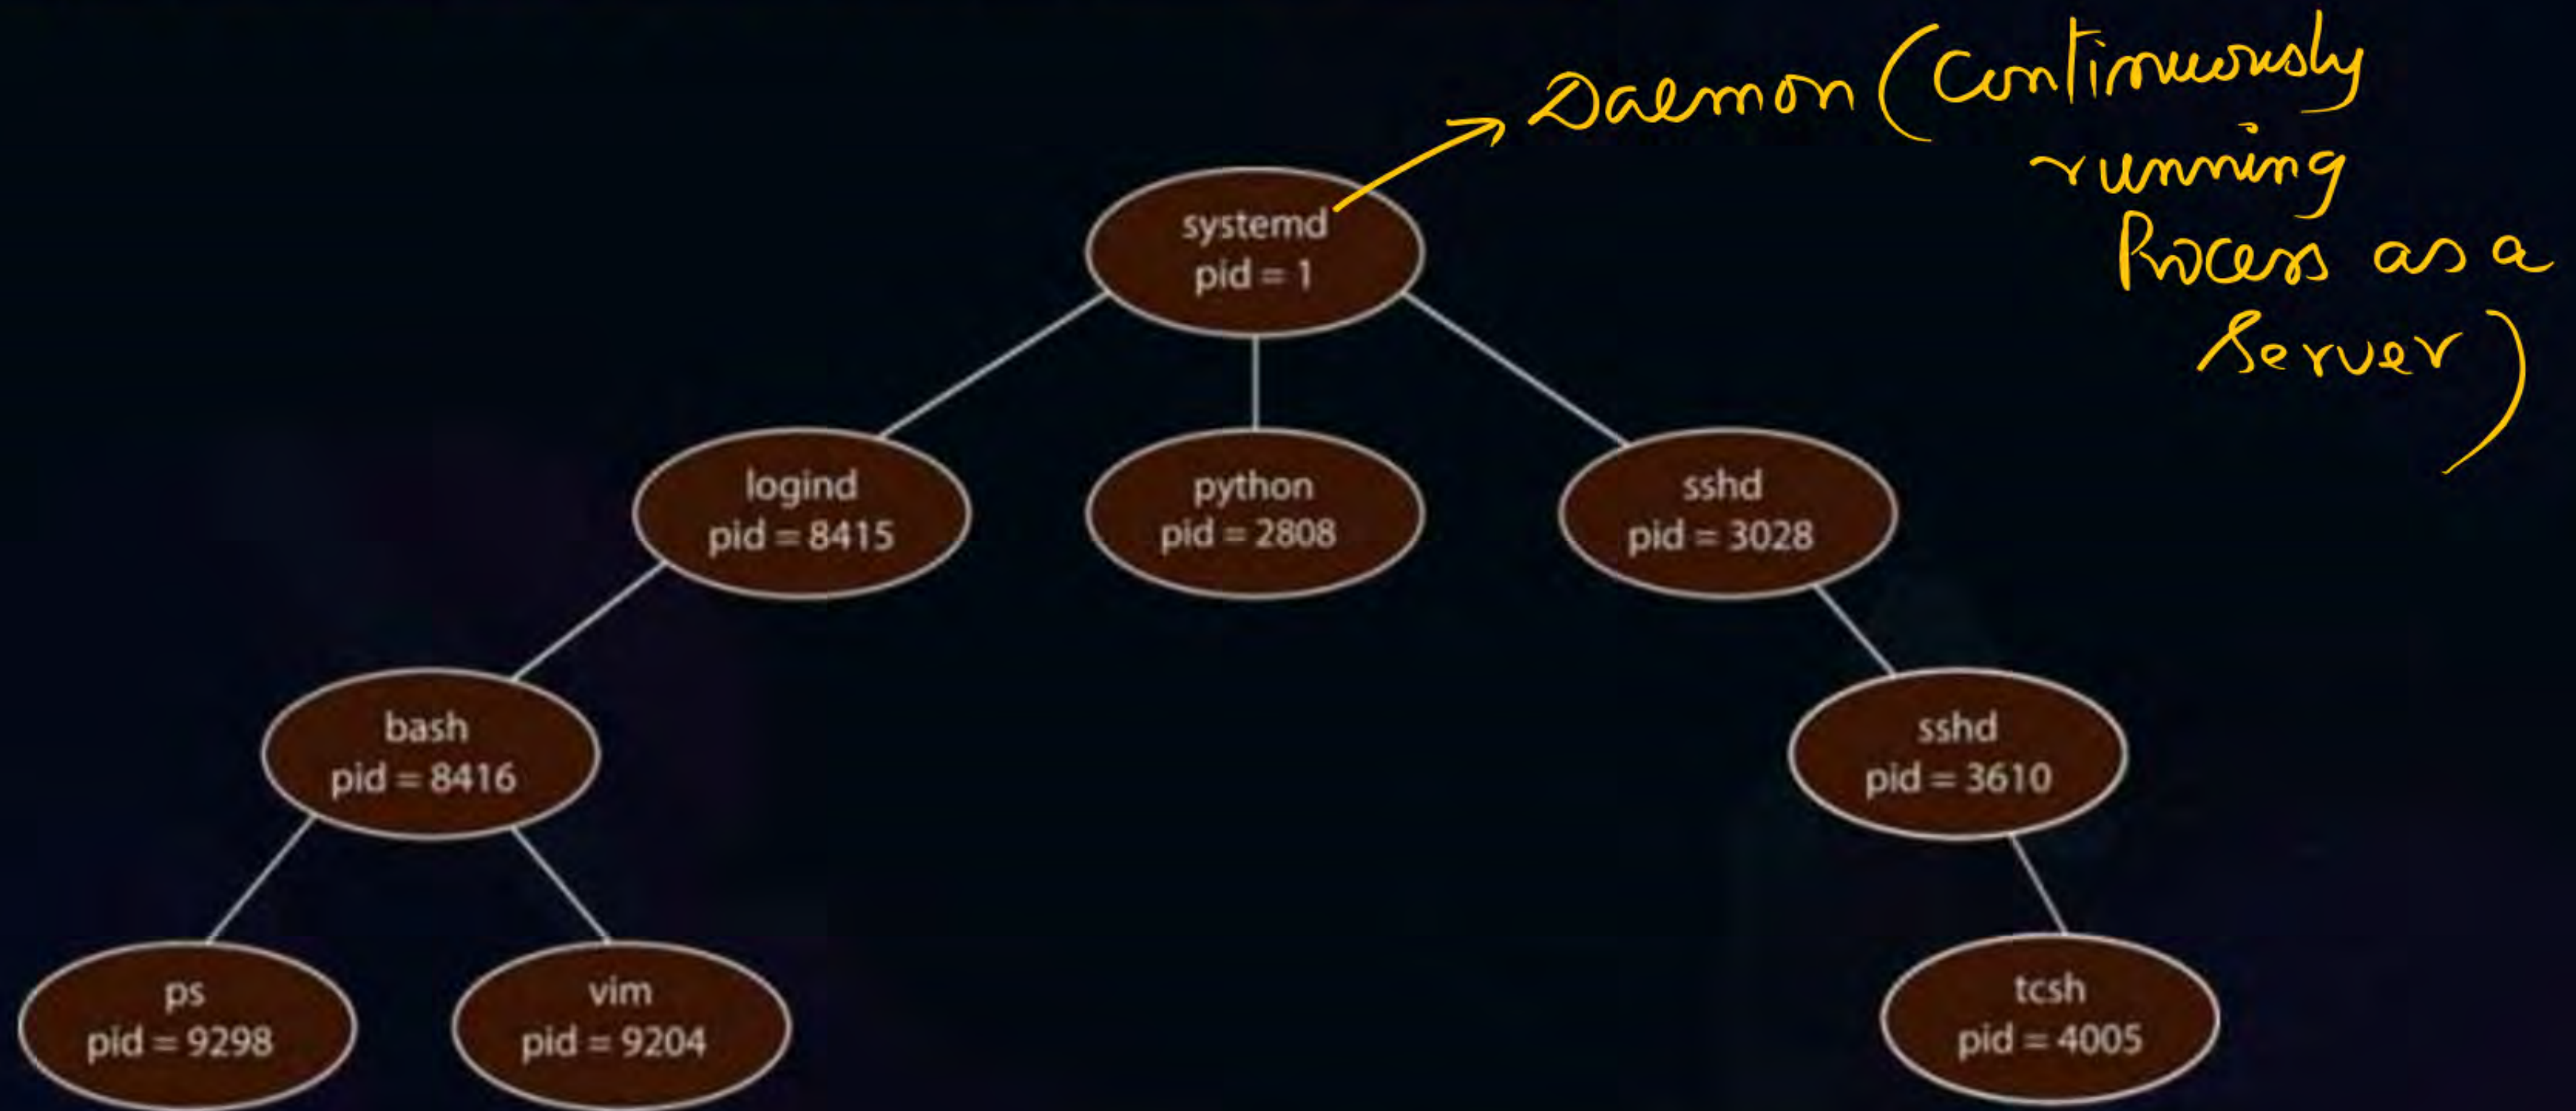
Topic : Process Creation (Cont.)

- Address space
- Child duplicate of parent
- Child has a program loaded into it
- UNIX examples
- ✓ ■ `fork()` system call creates new process
- ✓ ■ `exec()` system call used after a `fork()` to replace the process' memory space with a new program
- Parent process calls `wait()` waiting for the child to terminate





Topic : A Tree of Processes in Linux






Topic : C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{ pid_t pid;
  /* fork a child process */
  pid = fork() ;

  if (pid < 0) { /* error occurred */
    Trivial { fprintf(stderr, "Fork Failed");
              return 1;
            }
  }
  else if (pid == 0) { /* child process */
    execvp(" /bin/ls", "ls", NULL);
  }
```

```
}
else { /* parent process */
  /* parent will wait for the child to
  complete */
   $\Rightarrow$  (wait (NULL)); Parent waits for child to complete
  printf("Child Complete");
}
return 0;
}
```





Topic : Process Termination

- Process executes last statement and then asks the operating system to delete it using the `exit()` system call.
- Returns status data from child to parent (via `wait()`)
- Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the `abort()` system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting, and the operating systems does not allow a child to continue if its parent terminates



Topic : Process Termination

- Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
- cascading termination. All children, grandchildren, etc., are terminated.
- The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the `wait()` system call. The call returns status information and the pid of the terminated process
- `pid = wait(&status);`
- If no parent waiting (did not invoke `wait()`) process is a zombie
- If parent terminated without invoking `wait()`, process is an orphan




Topic : Threads & Multithreading



- Most modern applications are multithreaded
- Threads run within application
- Multiple tasks with the application can be implemented by separate threads
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- Process creation is heavy-weight while thread creation is light-weight
- Can simplify code, increase efficiency
- Kernels are generally multithreaded

Multi Core Architectures

i3; i5; i7; -

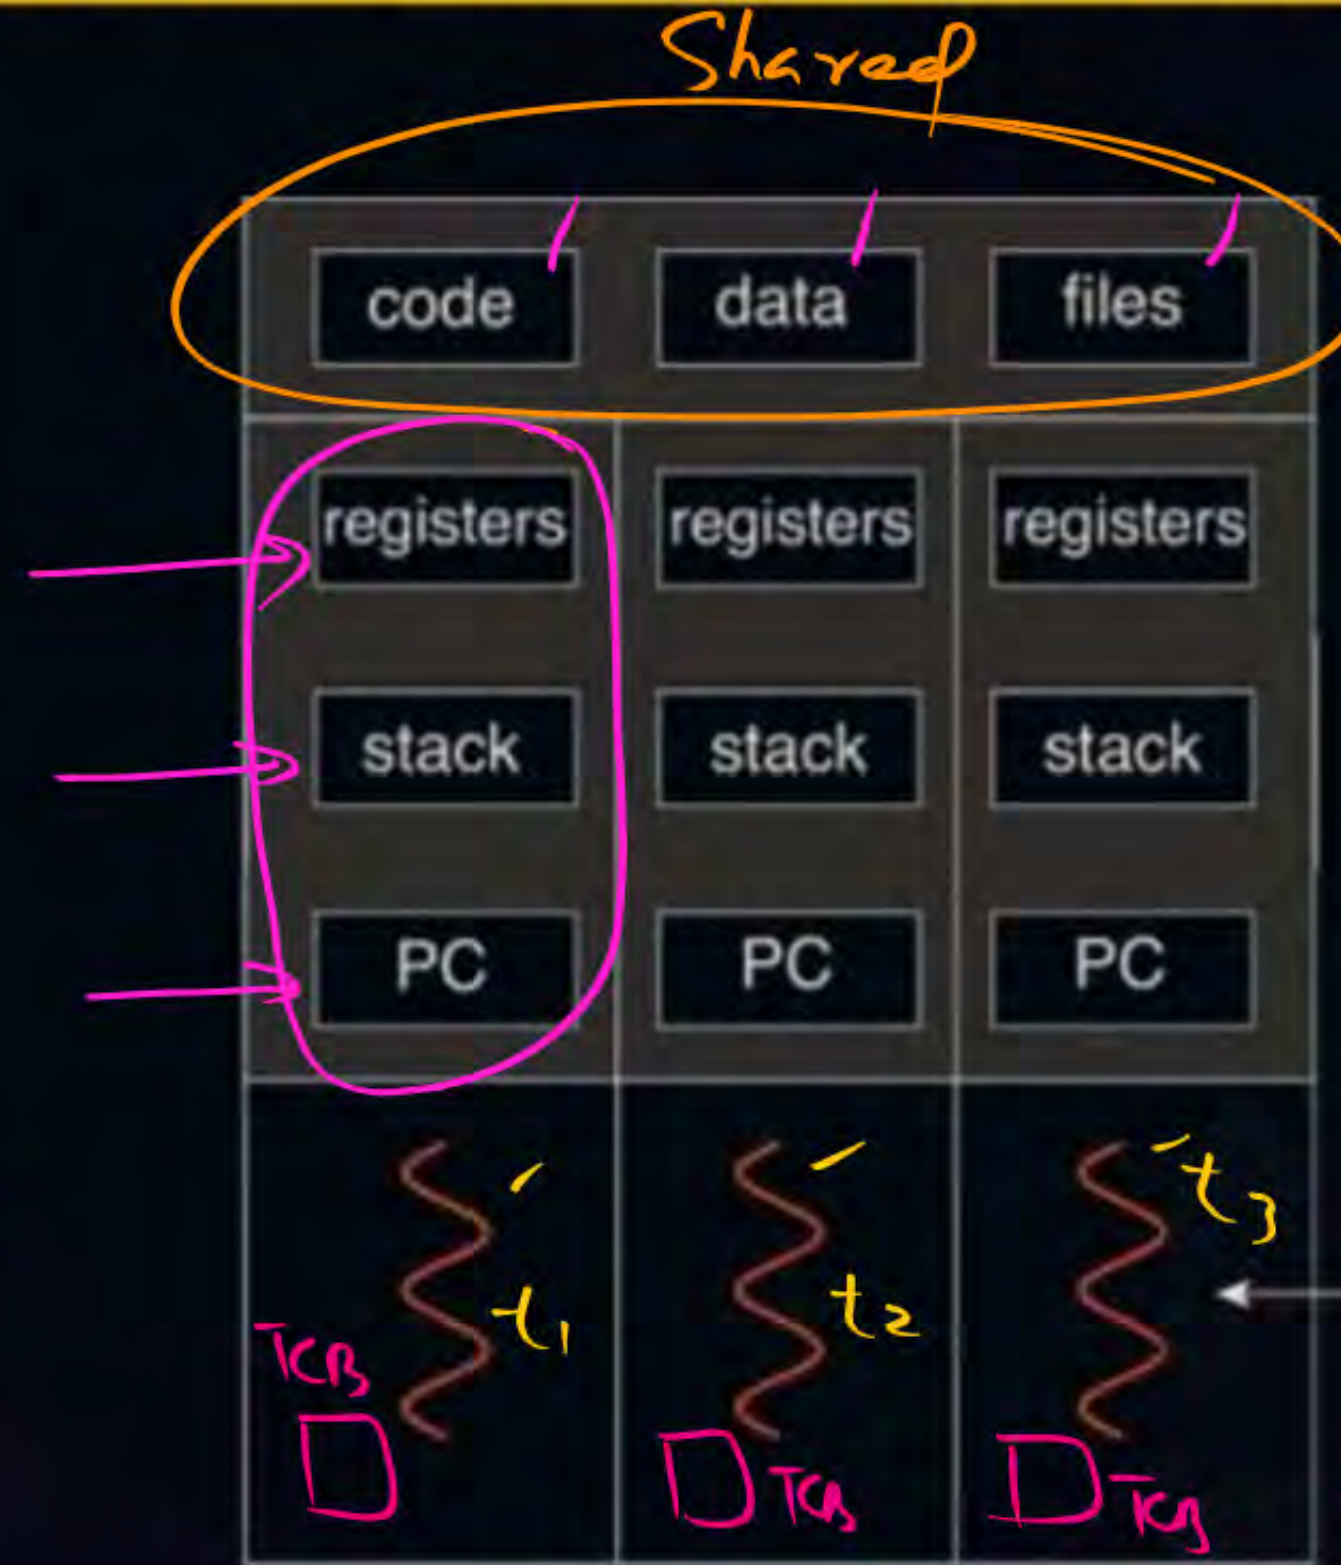




Topic : Single and Multithreaded Processes



single-threaded process



multithreaded process

→ Thread is a L.W.P
→ Part of Process

SSS

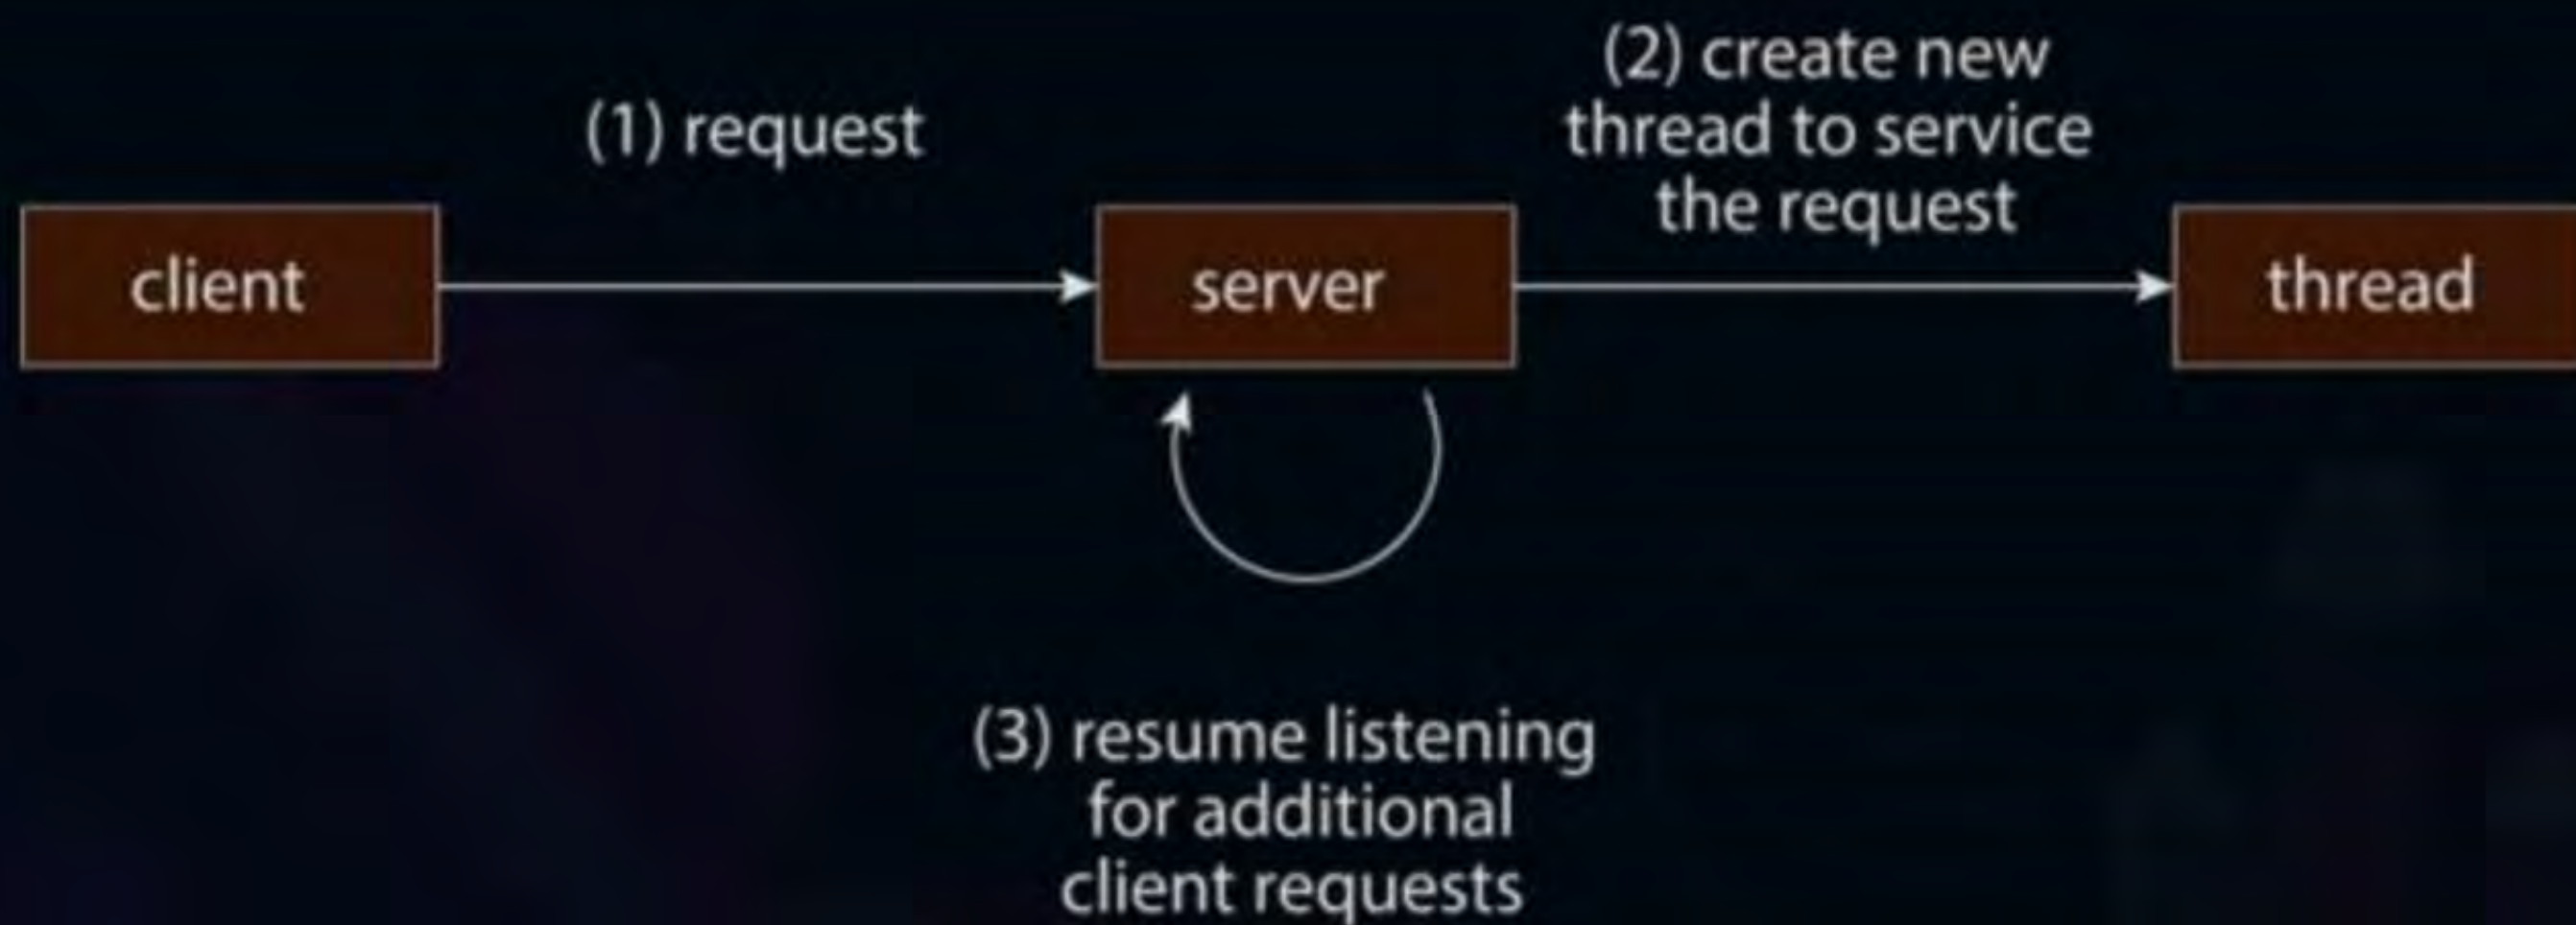
→ unit of CPU utilization

→ Thread can have its own Content (TCB)





Topic : Multithreaded Server Architecture





Topic : Benefits



- Responsiveness – may allow continued execution if part of process is blocked, especially important for user interfaces
- Resource Sharing – threads share resources of process, easier than shared memory or message passing
- Economy – cheaper than process creation, thread switching lower overhead than context switching
- Scalability – process can take advantage of multicore architectures.



Topic : Multicore Programming

- Multicore or multiprocessor systems puts pressure on programmers, challenges include:
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency



Topic : Concurrency vs. Parallelism

- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**



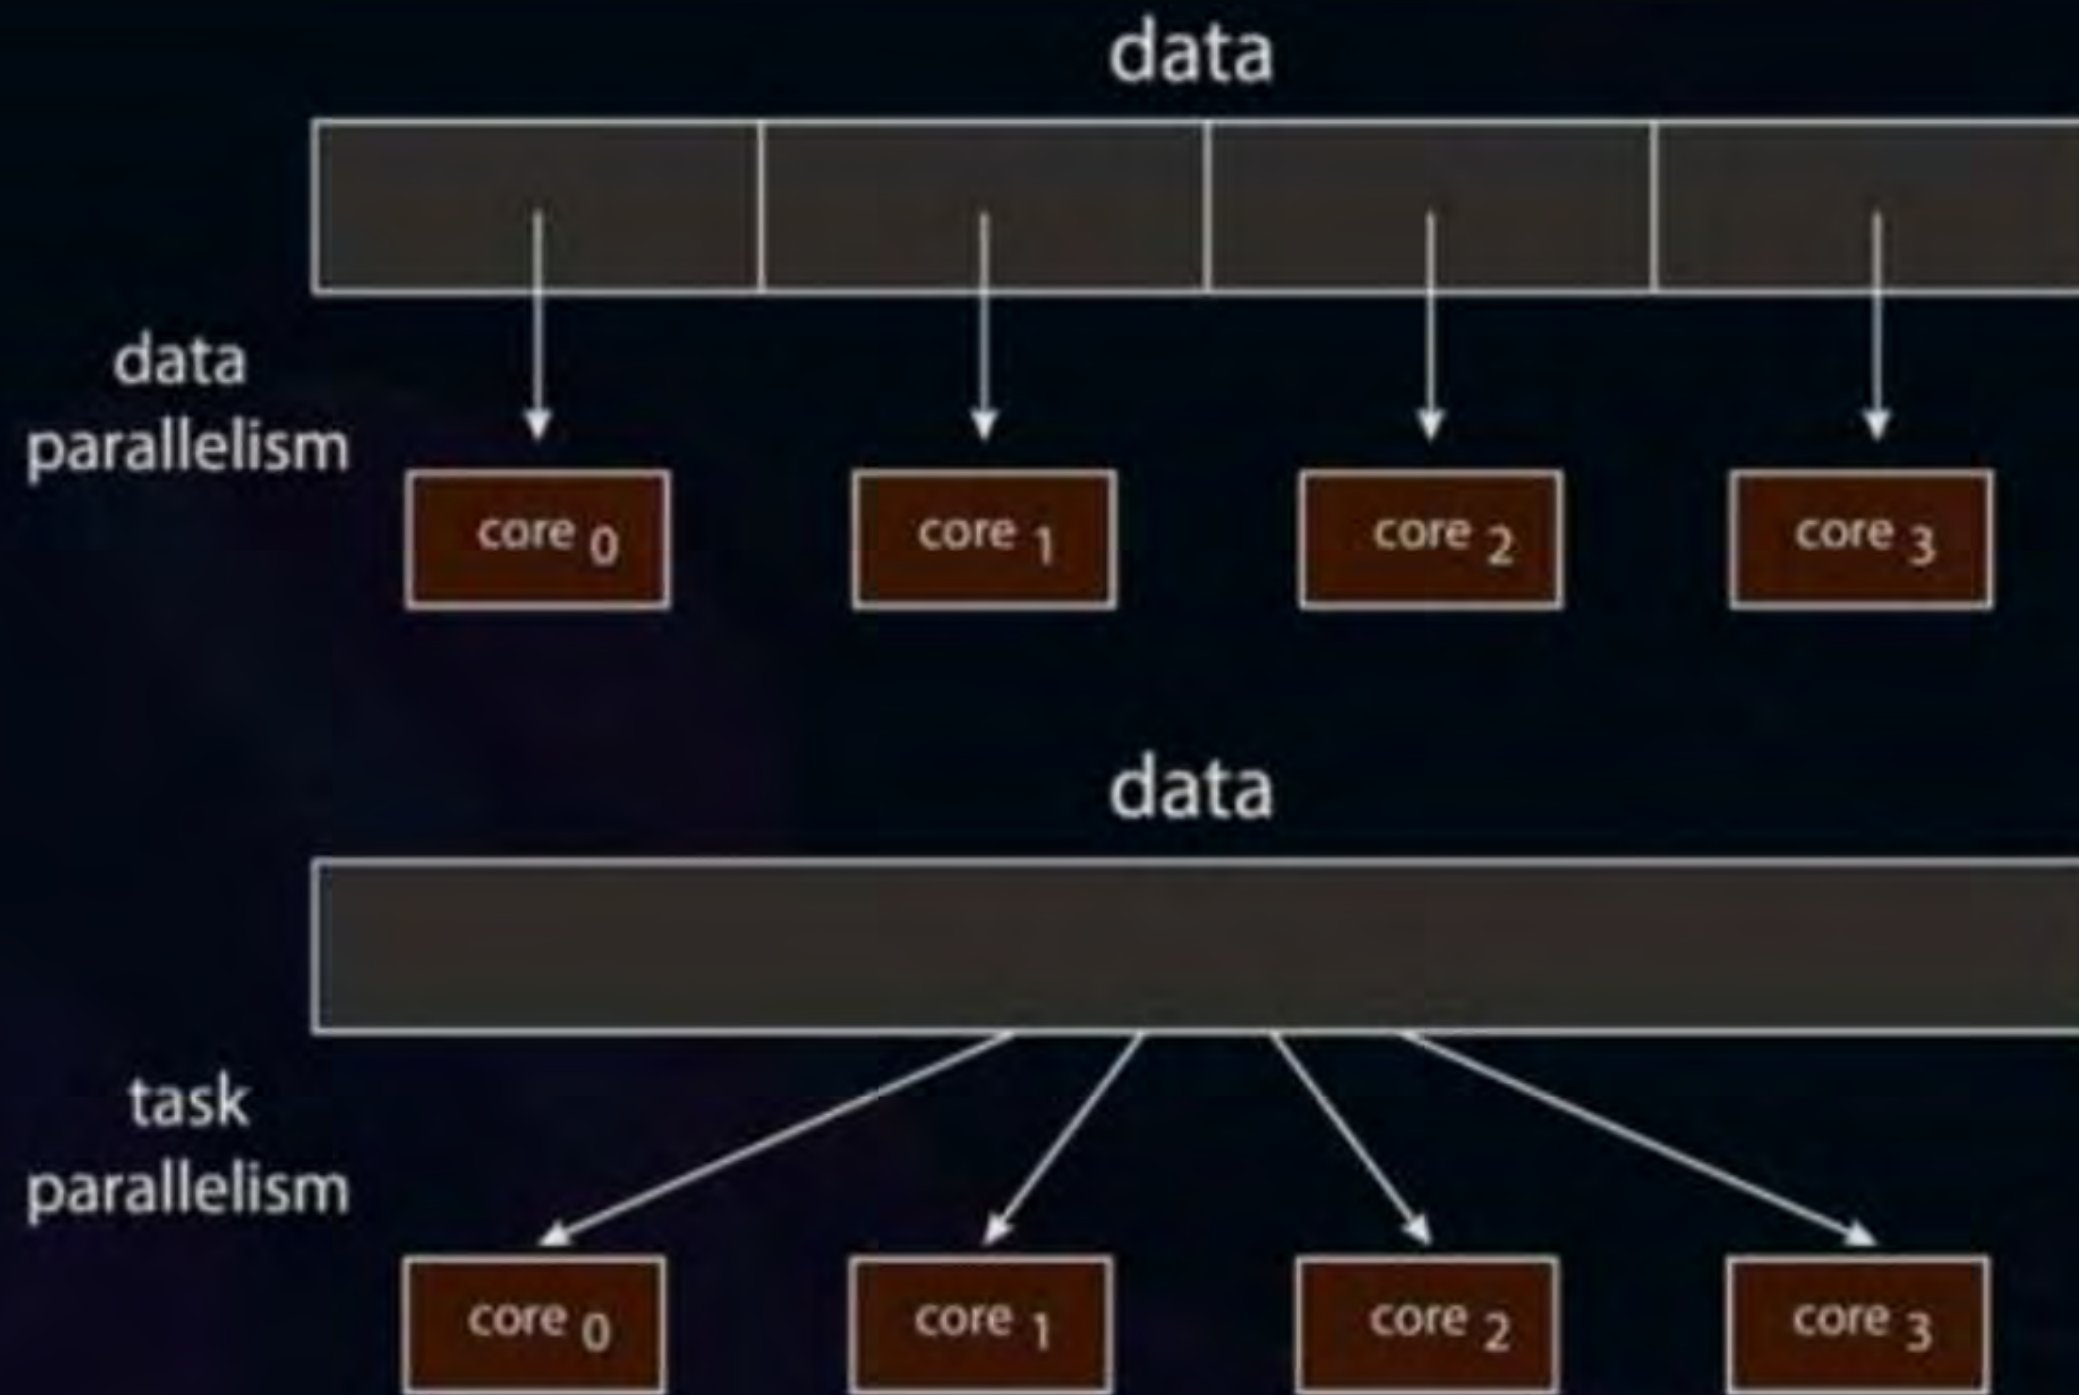


Topic : Multicore Programming

- Types of parallelism
 - Data parallelism – distributes subsets of the same data across multiple cores, same operation on each
 - Task parallelism – distributing threads across cores, each thread performing unique operation



Topic : Data and Task Parallelism





Topic : User Threads and Kernel Threads

Benefits of



U.L.T's

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX Pthreads ✓
 - Windows threads ✓
 - Java threads ✓
- **Kernel threads** - Supported by the Kernel
- Examples – virtually all general-purpose operating systems, including:
 - Windows
 - Linux
 - Mac OS X
 - iOS
 - Android


- 1) Flexibility
- 2) Transparency
- 3) Fastest CS
- 4) Improved Perf.

Drawback : of ULT

→ requirement of IO;
Sys Call by one ULT
will result in blocking
of whole process




Topic : User and Kernel Threads





Topic : Multithreading Models


- Many-to-One ✓
- One-to-One ✓
- Many-to-Many ✓





Topic : Many-to-One


- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads





Topic : One-to-One


- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux





Topic : Many-to-Many Model


- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the ThreadFiber package
- Otherwise not very common





Topic : Two-level Model

- Similar to M:M, except that it allows a user thread to be bound to kernel thread





Topic : Thread Libraries

- Thread library provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS



Topic : Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- Specification, not implementation
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)



Topic : Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */
void *runner (void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);
    /* wait for the thread to exit */
```




Topic : Pthreads Example (Cont.)

```
pthread_join (tid, NULL);

printf("sum = %d\n",sum);
}
/* The thread will execute in this function */
void *runner (void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```




Topic : Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
- Java threads may be created by:
 - Extending Thread class
 - Implementing the Runnable interface

```
Public interface Runnable  
{  
    public abstract void run();  
}
```

- The Executor is used as follows:



Topic : Java Threads

- **Implementing Runnable interface:**

```
class Task implements Runnable
{
    public void run() {
        System.out.println("I am a thread.");
    }
}
```

- **Creating a thread:**

```
Thread worker = new Thread (new Task());
Worker.start();
```

- **Waiting on a thread:**

```
Try {
    Worker.join();
}
Catch (InterruptedException ie.)
```




Topic : Java Executor Framework

- Rather than explicitly creating threads, Java also allows thread creation around the Executor interface:

```
public interface Executor
```

```
{
```

```
    void execute (Runnable command);
```

```
}
```

- The Executor is used as follows:

```
Executor service = new Executor;
```

```
Service.execute(new Task());
```




Topic : Java Executor Framework

```
import java.util.concurrent.*;

class Summation implements Callable<Integer> {
    private int upper;
    public Summation (int upper) {
        this.upper = upper;
    }
    /* The thread will execute in this method */ public Integer call() {
        int sum = 0;
        for (int i 1; i <= upper; i++) =
            sum += i;

        return new Integer (sum);
    }
}
```




Topic : Java Executor Framework (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        int upper = Integer.parseInt(args[0]);

        ExecutorService pool = Executors.newSingleThreadExecutor();
        Future<Integer> result = pool.submit(new Summation
(upper));

        try {
            System.out.println("sum " + result.get());
        } catch (InterruptedException | ExecutionException ie){}
    }
}
```


Q) Consider a System having n -cpu's & ' K ' Processes,
($K < n$);

Calculate the Max(Ready; Running; Block)

(1-CPU & ' K ' processes)

Ready; Running; Block
(K 1 K)



THANK - YOU